

Recreating Historical Chess Style with Policy-Only Fine-Tuning and Quantitative Validation

Nick Lamb

Medcopywriter Ltd, United Kingdom · ORCID 0009-0009-6266-8499

July 2026 · Preprint · CC BY 4.0 · DOI: 10.5281/zenodo.21564620

Interactive demonstration: chess.pharmatools.ai · Code (MIT): github.com/nickjlamb/time-machine-chess

Abstract

Chess engines optimize for move quality and, in doing so, converge on a single contemporary style. Chess *style*, however, is historically situated: opening repertoires, draw culture, and resignation habits all changed measurably between 1840 and 2019. We present Time-Machine Chess, a family of five era-specific chess models produced by one epoch of fine-tuning Maia-2 — a human-move-prediction network — on balanced corpora of dated over-the-board games (1840–1885, 1900–1939, 1950–1985, 1990–1999, 2010–2019). Models are served policy-only (no search), with a temperature schedule that preserves period opening diversity, and a rule-based “social layer” that derives era-appropriate draw agreements and resignations from the model’s own win-probability head. We validate the era claim three ways: (1) self-play reproduces historical opening, draw-rate, and game-length gradients within ~3 percentage points and ~5 plies respectively; (2) used as a generative classifier on held-out historical games, the models identify a game’s era at 42% (single game, 20% chance) and 88% (20-game batches), with a confusion structure that itself tracks chess history; and (3) strength measured against a calibrated engine ladder places all five models within a 181-point Elo band (~1580–1760), supporting the design goal that eras differ in style rather than strength. Total training cost is under one laptop-hour; serving requires no GPU.

1 Introduction

Superhuman chess engines are stylistically ahistorical: whatever their architecture, optimization pressure drives them toward the same narrow set of moves. Yet the human game has a documented stylistic history. The King’s Gambit appeared in 14% of surviving 1840–1885 master games and is effectively extinct in modern over-the-board play; draw rates rose from 12% to over 31% across 150 years before plateauing in the engine age; entire opening families rose and fell with theory and culture.

This work asks whether that history can be recreated as a playable artifact: a set of opponents that make the moves — and the non-move decisions, such as when to agree a draw or resign — characteristic of a specific era. The contribution is less any single technique than the combination of (a) a deliberately minimal modeling recipe (brief policy-only fine-tuning of an existing human-imitation model) with (b) a quantitative validation programme that treats every stylistic claim as testable against the historical record, and reports its residuals and noise floors.

2 Background

Maia [1] demonstrated that a policy network trained on human games predicts human moves far better than engines do, and that skill-conditioned variants capture rating-specific behavior. Maia-2 [2] unified this into a single model conditioned on the ratings of both players, trained on large volumes of online chess. Both model *what humans play* rather than what is best, which makes them natural substrates for style transfer: an era corpus is, in effect, a description of what a particular population of humans played.

We are not aware of prior work that fine-tunes human-imitation chess models on historical era corpora with quantitative validation against those corpora, though the components — fine-tuning, generative classification, engine-ladder strength estimation — are all standard.

3 Method

3.1 Data

Games come from Lumbra’s Gigabase [3], a database of ~10M dated over-the-board games (used under CC BY-NC-SA terms; not redistributed). A streaming filter partitions games by year into five era corpora: Romantic (1840–1885; 10,693 games — effectively the entire surviving record), Classical (1900–1939; 62,779), Soviet (1950–1985; 597,511), Engine Dawn (1990–1999; 500,000), and Engine Era (2010–2019; 500,000). Pre-1840 chess is data-starved — roughly 670 game scores survive — and was excluded.

From each corpus we draw a balanced fine-tuning subset (10,000–12,000 games; 0.73M to 0.95M positions per era) so that eras receive comparable training signal despite corpus sizes varying by a factor of 56.

3.2 Fine-tuning

Each era model is initialized from the pretrained Maia-2 rapid checkpoint and fine-tuned for a single epoch on its era subset (Apple M-series laptop, minutes per era). Positions with Black to move are mirrored per Maia-2 convention. Because historical games generally lack ratings, all positions are conditioned on a fixed nominal rating of 1900 for both players, at training and at inference. The intent is that era models should differ in *style* while holding strength approximately constant; Section 4.3 tests whether this held.

3.3 Policy-only serving

At play time the era model’s policy head yields a probability distribution over legal moves, from which the served move is sampled directly — no search of any kind. The first 10 plies are sampled at temperature 1.0, later moves at temperature 0.6: era character is disproportionately carried by opening *diversity*, which sharpening destroys, while unsharpened middlegame sampling produces uncharacteristic errors. Serving is CPU-only (~50–100 ms per move; checkpoints are 89 MB).

3.4 The social layer

Draw agreements and resignations are not moves, and a policy network therefore never learns them; without them, initial validation showed games ~25 plies too long and draw rates 8–10 points below history. We model both as era-specific threshold rules on the model’s own win-probability head: a draw is offered/agreed when the evaluation has remained within a dead-equal band for a required number of consecutive moves past a minimum move number; the model resigns when its own win probability has remained below a threshold for a required number of moves. Constants were tuned per era against the historical record. The two rules interact — prompt resignation consumes games that would otherwise have reached draw-agreement territory — and were tuned jointly.

4 Validation

4.1 Self-play against the historical record

Each era model played 150 self-play games at serving temperature; identical move-sequence metrics (no annotation metadata) were computed on these games and on random samples of the corresponding era corpus. Table 1 reports the headline metrics as historical / bot values.

Metric	Romantic	Classical	Soviet	E. Dawn	E. Era
King's Gambit %	14.0 / 18.0	3.5 / 9.3	1.25 / 0.0	0.75 / 0.0	0.0 / 0.0
Draw rate %	12.0 / 15.3	25.0 / 24.7	28.75 / 26.0	31.75 / 34.7	31.75 / 28.7
Mean length (plies)	73.9 / 79.1	77.5 / 74.5	72.7 / 69.1	75.2 / 70.8	78.4 / 78.0

Table 1: Self-play vs. historical corpus (historical / bot).

The historical gradients reproduce: the King’s Gambit dies on schedule; draw rates rise for 150 years and plateau in the engine age (the plateau itself — coincident with games reaching their historical maximum length — is a finding the tuning did not target); the 2010s model plays 1.c4 at exactly the historical 8.0%. The principal residual is a first-move lean toward 1.e4 beyond historical rates in older eras: one epoch moves Maia-2’s modern opening prior most, but not all, of the way. At 150 games the draw-rate standard error is ~ 3.5 points; social-layer constants were tuned to within ~ 3 points of history and no further.

4.2 Held-out era classification

If each model best predicts its own era’s moves, the ensemble is a generative classifier, and its accuracy on held-out games is an independent test of era separability. For each game, each era model assigns a probability to every move actually played (positions sampled after the first 6 plies, excluding forced moves and decided positions); the era with the highest summed log-likelihood is the prediction. Test games (100 per era) were drawn from the corpora with the training subsets reproduced by seed and excluded.

true \ predicted	Rom.	Cla.	Sov.	Dawn	Era
Romantic	60	14	2	12	12
Classical	15	42	12	16	15
Soviet	17	13	28	22	20
E. Dawn	10	15	12	43	20
E. Era	14	4	10	35	37

Table 2: Confusion matrix, single held-out games ($n = 100$ per era; chance = 20%).

Single-game accuracy is 42% against a 20% chance baseline; aggregating evidence over 20-game batches yields 88%. The error structure is historically interpretable: the Romantic era is by far the most identifiable; the two engine-age eras substantially merge (scored as one class, their combined accuracy is $\sim 68\%$), consistent with a converged modern game; and the Soviet era is the weakest diagonal, its sound, theory-driven play passing plausibly in any later era.

4.3 Measured playing strength

The fixed rating conditioning (Section 3.2) is an input, not a measurement. Each era model therefore played 120 games against Stockfish 16 [4] at four limited-strength anchor levels (UCI_Elo 1400/1700/2000/2300; colors alternated; era resignation rules applied), and a performance rating was fitted by maximum likelihood with a bootstrap 95% interval.

Era	Fitted rating	95% CI
Romantic	1690	1597–1786
Classical	1579	1487–1677
Soviet	1652	1561–1742
Engine Dawn	1760	1656–1862
Engine Era	1701	1607–1794

Table 3: Engine-referenced strength estimates (120 games per era).

All five models fall within a 181-point band with overlapping intervals — roughly one rating class — supporting the style-not-strength design goal to first order. Stockfish’s limited-strength mode itself approximates human rating, and engine-calibrated ratings transfer imperfectly to human opponents; we report these as engine-referenced estimates.

5 Limitations

The models inherit Maia-2’s modern online-chess prior, visible as the residual 1.e4 lean. Policy-only serving means no tactical verification: in positions far from any era’s data the move distribution flattens, and play degrades — a property consistent with the goal of human-like era play but unsuitable for applications requiring soundness. The social layer is rule-based, not learned; its constants are tuned to aggregate era statistics and interact with each other. Era classification is noisy at single-game granularity by the nature of the signal, and adjacent eras are genuinely non-separable in part. Strength estimates are engine-calibrated. Finally, the historical record itself is a biased sample — early corpora preserve disproportionately notable games — and “era style” as learned here is the style of *recorded* games.

6 Conclusion

A minimal recipe — brief policy-only fine-tuning of a human-imitation model on era corpora, plus a rule-based social layer driven by the model’s own evaluation — suffices to recreate measurable, validated historical chess style at negligible compute cost. The validation programme was not overhead on the artifact: it produced the principal findings, including the draw-culture plateau, the near-merger of the engine-age eras, the Soviet school’s stylistic universality, and the gap between conditioning inputs and measured strength. We suggest the same imitate/fine-tune/validate pattern applies to recreating historical or school styles in other domains where dated human decision records exist.

Data and code availability

Source code, validation scripts, and all validation outputs are available under MIT license at github.com/nickjlamb/time-machine-chess. Model checkpoints are distributed via the repository’s releases. Training data (Lumbra’s Gigabase [3]) is CC BY-NC-SA and is not redistributed. An interactive demonstration, including the era classifier and all validation results, is at chess.pharmatools.ai.

Acknowledgements

Built on Maia-2 by the Computational Social Science Lab, University of Toronto [2]. Historical games from Lumbra’s Gigabase [3]. Board interaction uses `python-chess` [5].

References

- [1] R. McIlroy-Young, S. Sen, J. Kleinberg, A. Anderson. Aligning Superhuman AI with Human Behavior: Chess as a Model System. *Proc. 26th ACM SIGKDD*, 2020. arXiv:2006.01855.

- [2] Z. Tang, D. Jiao, R. McIlroy-Young, J. Kleinberg, S. Sen, A. Anderson. Maia-2: A Unified Model for Human-AI Alignment in Chess. *NeurIPS*, 2024. arXiv:2409.20553.
- [3] Lumbr’s Gigabase. lumbrasgigabase.com (accessed July 2026).
- [4] The Stockfish developers. Stockfish 16. stockfishchess.org.
- [5] N. Fiekas. python-chess. github.com/niklasf/python-chess.
- [6] A. Elo. *The Rating of Chessplayers, Past and Present*. Arco, 1978.